

## Practical – 3

**Aim:** Implement and demonstrate the FIND-S Algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a .CSV file

- **Code:**

```
import pandas as pd
df = pd.read_csv("/ws.csv")
print("Columns in ws.csv:", df.columns)
features = df.columns[:-1]
target = df.columns[-1]
df[features] = df[features].astype(str)
df[target] = df[target].astype(str)
positive_example = df[df[target] == 'Yes'][features].values
def find_s(example):
    if len(example) == 0:
        return "No positive examples found."
    hypothesis = example[0].copy()
    print("Initial hypothesis:", hypothesis)
    for i, instance in enumerate(example[1: ]):
        for j in range(len(hypothesis)):
            if hypothesis[j] != instance[j]:
                hypothesis[j] = '?'
        print(f'Hypothesis after example {i+2}:', hypothesis)
    return hypothesis
hypothesis = find_s(positive_example)
print("Most specific hypothesis:", hypothesis)
```

- **Output:**

```
Columns in ws.csv: Index(['Sunny', 'Warm', 'Normal', 'Strong', 'Warm.1', 'Same', 'Yes'], dtype='object')
Initial hypothesis: ['Sunny' 'Warm' 'High' 'Strong' 'Warm' 'Same']
Hypothesis after example 2: ['Sunny' 'Warm' 'High' 'Strong' '?' '?']
Most specific hypothesis: ['Sunny' 'Warm' 'High' 'Strong' '?' '?']
```

## Practical – 4

**Aim:** Write a Python program to implement Simple Linear Regression

- How many total observations in data?
- How many independent variables?
- Which is a dependent variable?
- Quantify the goodness of your model and discuss steps taken for improvement (RMSE, SSE, R2Score).

- **Code:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from pandas.core.common import random_state
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

df = pd.read_csv("/Salary_Data.csv")
df.head()
df.describe()
X = df.iloc[:, :1] # independent
y = df.iloc[:, 1:] # dependent
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
regressor = LinearRegression()
regressor.fit(X_train, y_train)
y_pred_test = regressor.predict(X_test)
y_pred_train = regressor.predict(X_train)
rmse = np.sqrt(mean_squared_error(y_test, y_pred_test))
sse = np.sum((y_test - y_pred_test) ** 2)
r2 = r2_score(y_test, y_pred_test)
print(f"\nModel Evaluation:")
print(f"RMSE: {rmse:.4f}")
print(f"SSE: {sse.iloc[0]:.4f}")
print(f"R2 Score: {r2:.4f}")
plt.scatter(X_train, y_train, color = 'green')
plt.plot(X_train, y_pred_train, color = 'firebrick')
plt.title('Salary vs Experience (Training Set)')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.legend(['X_train/Pred(y_test)', 'X_train/y_train'], title = 'Sal/Exp', loc='best',
facecolor='lightblue')
plt.box(False)
plt.show()
```

- **Output:**

Model Evaluation:  
RMSE: 3580.9792  
SSE: 76940473.7888  
R<sup>2</sup> Score: 0.9882



## Practical – 5

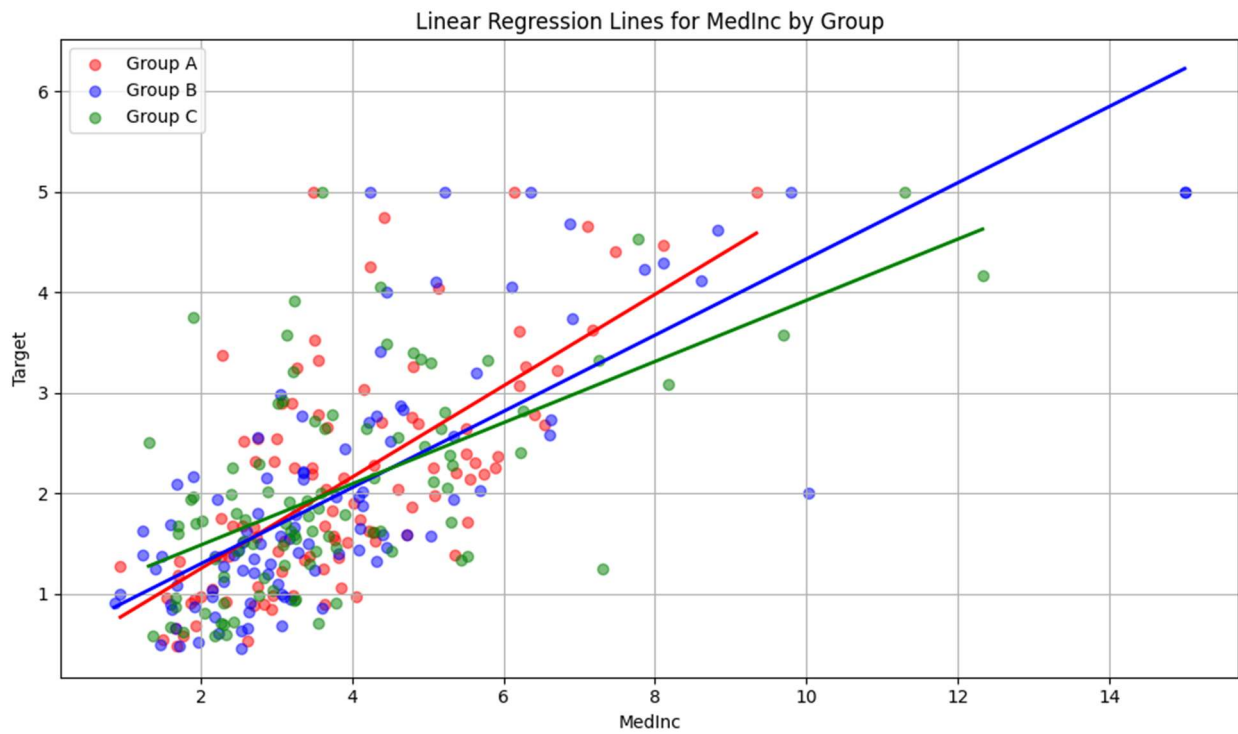
**Aim:** Implementation of Multiple Linear Regression for House Price Prediction using sklearn.

- **Code:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
df = pd.DataFrame(housing.data, columns=housing.feature_names)
df['Target'] = housing.target
df = df.sample(300, random_state=42).reset_index(drop=True)
df['Group'] = np.random.choice(['A', 'B', 'C'], size=len(df))
x_feature = 'MedInc'
y_feature = 'Target'
plt.figure(figsize=(10, 6))
colors = {'A': 'red', 'B': 'blue', 'C': 'green'}
for group in df['Group'].unique():
    subset = df[df['Group'] == group]
    X = pd.DataFrame(subset[[x_feature]])
    y = subset[y_feature]
    model = LinearRegression()
    model.fit(X, y)
    x_line = np.linspace(X[x_feature].min(), X[x_feature].max(), 100).reshape(-1, 1)
    y_line = model.predict(pd.DataFrame(x_line, columns=[x_feature]))
    plt.scatter(X, y, alpha=0.5, label=f'Group {group}', color=colors[group])
    plt.plot(x_line, y_line, color=colors[group], linewidth=2)
plt.xlabel(x_feature)
plt.ylabel(y_feature)
plt.title(f'Linear Regression Lines for {x_feature} by Group')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

- **Output:**



## Practical – 6

### Aim: Two Class Classification (Logistic Regression)

- How many total observations in data?
- How many independent variables?
- Which is a dependent variable?
- Implement logistic function.
- Implement Log-loss function.
- Quantify the goodness of your model and discuss steps taken for improvement (Accuracy, Confusion matrices, F-measure).

- **Code:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
    accuracy_score, confusion_matrix, f1_score, log_loss, classification_report
)
housing = fetch_california_housing()
df = pd.DataFrame(housing.data, columns=housing.feature_names)
df['Target'] = housing.target
median_value = df['Target'].median()
df['Target_Binary'] = (df['Target'] > median_value).astype(int)
X = df.drop(columns=['Target', 'Target_Binary'])
y = df['Target_Binary']
print("Total observations:", len(df))
print("Number of independent variables:", X.shape[1])
print("Dependent variable: Target_Binary")
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
y_pred_prob = model.predict_proba(X_test)[:, 1]
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
logloss = log_loss(y_test, y_pred_prob)
print("\nAccuracy:", accuracy)
print("F1 Score:", f1)
print("Log Loss (sklearn):", logloss)
print("\nClassification Report:\n", classification_report(y_test, y_pred))
def logistic(z):
    return 1 / (1 + np.exp(-z))
def log_loss_manual(y_true, y_prob):
```

```

eps = 1e-15
y_prob = np.clip(y_prob, eps, 1 - eps)
return -np.mean(y_true * np.log(y_prob) + (1 - y_true) * np.log(1 - y_prob))
manual_logloss = log_loss_manual(y_test.values, y_pred_prob)
print("Manual Log Loss:", manual_logloss)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()

```

- **Output:**

```

Total observations: 20640
Number of independent variables: 8
Dependent variable: Target_Binary

Accuracy: 0.8273578811369509
F1 Score: 0.8281074127673259
Log Loss (sklearn): 0.38525281617941143


Classification Report:
              precision    recall  f1-score   support

     0       0.82         0.83         0.83        3069
     1       0.83         0.82         0.83        3123

 accuracy          0.83         0.83         0.83        6192
 macro avg         0.83         0.83         0.83        6192
 weighted avg         0.83         0.83         0.83        6192

Manual Log Loss: 0.38525281617941143

```



## Practical – 7

**Aim: Implementation of Decision tree using sklearn and its parameter tuning.**

- **Code:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

housing = fetch_california_housing()
df = pd.DataFrame(housing.data, columns=housing.feature_names)
df['Target'] = housing.target
df = df.sample(n=200, random_state=42).reset_index(drop=True)
df['Target_Binary'] = (df['Target'] > df['Target'].median()).astype(int)
X = df.drop(columns=['Target', 'Target_Binary'])
y = df['Target_Binary']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
clf = DecisionTreeClassifier(random_state=42)
param_grid = {
    'max_depth': [3, 5, None],
    'min_samples_split': [2, 10],
    'criterion': ['gini', 'entropy']
}
grid = GridSearchCV(clf, param_grid, cv=5, scoring='accuracy')
grid.fit(X_train, y_train)
best_tree = grid.best_estimator_
print("Best Parameters:", grid.best_params_)
y_pred = best_tree.predict(X_test)
acc = accuracy_score(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)
print("\nAccuracy:", acc)
print("\nConfusion Matrix:\n", cm)
print("\nClassification Report:\n", classification_report(y_test, y_pred))
plt.figure(figsize=(16, 8))
plot_tree(best_tree, feature_names=X.columns, class_names=["Low", "High"], filled=True,
rounded=True)
plt.title("Decision Tree Visualization")
plt.show()
```



- **Output:**

Best Parameters: {'criterion': 'gini', 'max\_depth': 3, 'min\_samples\_split': 2}

Accuracy: 0.75

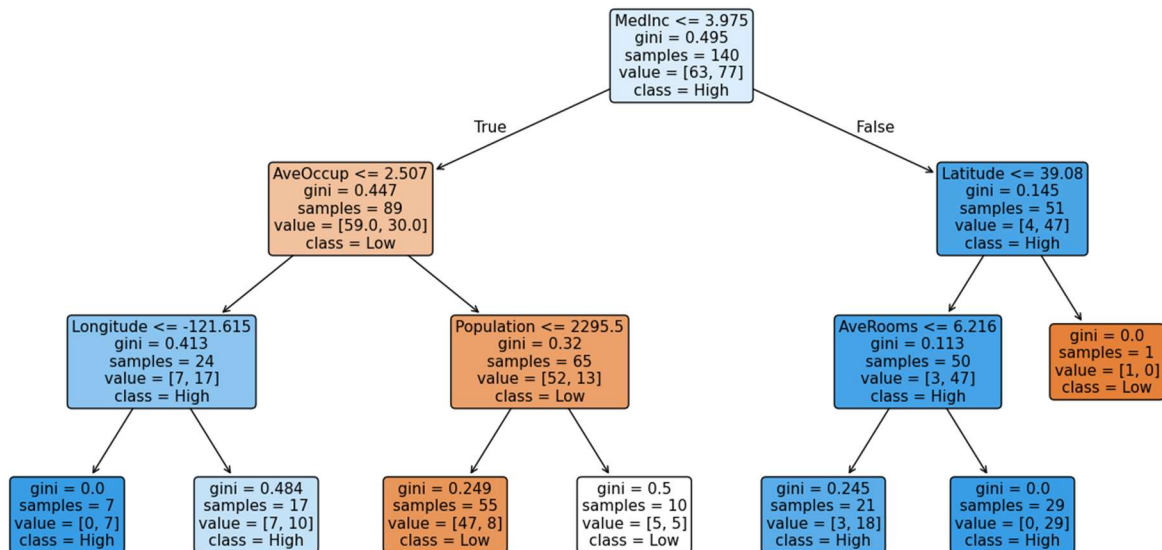
Confusion Matrix:

```
[[25 12]
 [ 3 20]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.68	0.77	37
1	0.62	0.87	0.73	23
accuracy			0.75	60
macro avg	0.76	0.77	0.75	60
weighted avg	0.79	0.75	0.75	60

Decision Tree Visualization



## Practical – 9

### Aim: Write a program to implement Random Forest Algorithm

- **Code:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_california_housing
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

housing = fetch_california_housing()
df = pd.DataFrame(housing.data, columns=housing.feature_names)
df['Target'] = housing.target
df = df.sample(n=300, random_state=42).reset_index(drop=True)
median_value = df['Target'].median()
df['Target_Binary'] = (df['Target'] > median_value).astype(int)
X = df.drop(columns=['Target', 'Target_Binary'])
y = df['Target_Binary']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
rf = RandomForestClassifier(
    n_estimators=100,
    max_depth=None,
    random_state=42
)
rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)
acc = accuracy_score(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)
print("Accuracy:", acc)
print("\nConfusion Matrix:\n", cm)
print("\nClassification Report:\n", classification_report(y_test, y_pred))
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()
importances = rf.feature_importances_
feature_names = X.columns
sorted_idx = np.argsort(importances)[::-1]
plt.figure(figsize=(10, 6))
sns.barplot(x=importances[sorted_idx], y=feature_names[sorted_idx])
plt.title("Feature Importances in Random Forest")
plt.tight_layout()
plt.show()
```



- **Output:**

Confusion Matrix:

```
[[40  7]
 [11 32]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.85	0.82	47
1	0.82	0.74	0.78	43
accuracy			0.80	90
macro avg	0.80	0.80	0.80	90
weighted avg	0.80	0.80	0.80	90



## Practical – 10

### Aim: Write a program to implement Random Forest Algorithm

- **Code:**

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv("/content/drive/MyDrive/temp/1_Housing.csv")
# converting data into binaries 1 for high and 0 for low
median_price = df['price'].median()
df['price_category'] = (df['price'] > median_price).astype(int)
df.drop(columns=['price'], inplace=True)
categorical_cols = df.select_dtypes(include='object').columns
le = LabelEncoder()
for col in categorical_cols:
    df[col] = le.fit_transform(df[col])
X = df.drop(columns=['price_category'])
y = df['price_category']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
model = GaussianNB()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)
print("Accuracy:", accuracy)
print("\nConfusion Matrix:\n", cm)
print("\nClassification Report:\n", classification_report(y_test, y_pred))
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Naïve Bayes Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()
```

- **Output:**




Accuracy: 0.823170731707317

Confusion Matrix:

```
[[70 16]
 [13 65]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.81	0.83	86
1	0.80	0.83	0.82	78
accuracy			0.82	164
macro avg	0.82	0.82	0.82	164
weighted avg	0.82	0.82	0.82	164



## Practical – 11

**Aim: Write a program to implement OR, AND gate using Perceptron with learning rules.**

- **Code:**

```
import numpy as np

def step_function(x):
    return 1 if x >= 0 else 0

def train_perceptron(inputs, labels, learning_rate=0.1, iterations=5):
    weights = np.zeros(inputs.shape[1])
    bias = 0
    for iteration in range(iterations):
        print(f"\nIteration no: {iteration+1}")
        for x, label in zip(inputs, labels):
            z = np.dot(x, weights) + bias
            y_pred = step_function(z)
            error = label - y_pred
            weights += learning_rate * error * x
            bias += learning_rate * error
        print(f"Input: {x}, Pred: {y_pred}, Error: {error}, Updated Weights: {weights}, Bias: {bias}")
    return weights, bias

def predict(inputs, weights, bias):
    return [step_function(np.dot(x, weights) + bias) for x in inputs]

inputs = np.array([[0,0],[0,1],[1,0],[1,1]])
or_labels = np.array([0,1,1,1])
and_labels = np.array([0,0,0,1])
# OR Gate
print("Training Perceptron for OR Gate")
or_weights, or_bias = train_perceptron(inputs, or_labels)
or_outputs = predict(inputs, or_weights, or_bias)
print("\nFinal OR Predictions:")
for i, o in zip(inputs, or_outputs):
    print(f"Input: {i}, Output: {o}")
# AND Gate
print("\n\nTraining Perceptron for AND Gate")
and_weights, and_bias = train_perceptron(inputs, and_labels)
and_outputs = predict(inputs, and_weights, and_bias)
print("\nFinal AND Predictions:")
for i, o in zip(inputs, and_outputs):
    print(f"Input: {i}, Output: {o}")
```

- **Output:**

#### Training Perceptron for OR Gate

```

Iteration no: 1
Input: [0 0], Pred: 1, Error: -1, Updated Weights: [0. 0.], Bias: -0.1
Input: [0 1], Pred: 0, Error: 1, Updated Weights: [0. 0.1], Bias: 0.0
Input: [1 0], Pred: 1, Error: 0, Updated Weights: [0. 0.1], Bias: 0.0
Input: [1 1], Pred: 1, Error: 0, Updated Weights: [0. 0.1], Bias: 0.0

Iteration no: 2
Input: [0 0], Pred: 1, Error: -1, Updated Weights: [0. 0.1], Bias: -0.1
Input: [0 1], Pred: 1, Error: 0, Updated Weights: [0. 0.1], Bias: -0.1
Input: [1 0], Pred: 0, Error: 1, Updated Weights: [0.1 0.1], Bias: 0.0
Input: [1 1], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: 0.0

Iteration no: 3
Input: [0 0], Pred: 1, Error: -1, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [0 1], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [1 0], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [1 1], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1

Iteration no: 4
Input: [0 0], Pred: 0, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [0 1], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [1 0], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [1 1], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1

Iteration no: 5
Input: [0 0], Pred: 0, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [0 1], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [1 0], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [1 1], Pred: 1, Error: 0, Updated Weights: [0.1 0.1], Bias: -0.1

Final OR Predictions:
Input: [0 0], Output: 0
Input: [0 1], Output: 1
Input: [1 0], Output: 1
Input: [1 1], Output: 1

```

#### Training Perceptron for AND Gate

```

Iteration no: 1
Input: [0 0], Pred: 1, Error: -1, Updated Weights: [0. 0.], Bias: -0.1
Input: [0 1], Pred: 0, Error: 0, Updated Weights: [0. 0.], Bias: -0.1
Input: [1 0], Pred: 0, Error: 0, Updated Weights: [0. 0.], Bias: -0.1
Input: [1 1], Pred: 0, Error: 1, Updated Weights: [0.1 0.1], Bias: 0.0

Iteration no: 2
Input: [0 0], Pred: 1, Error: -1, Updated Weights: [0.1 0.1], Bias: -0.1
Input: [0 1], Pred: 1, Error: -1, Updated Weights: [0.1 0. ], Bias: -0.2
Input: [1 0], Pred: 0, Error: 0, Updated Weights: [0.1 0. ], Bias: -0.2
Input: [1 1], Pred: 0, Error: 1, Updated Weights: [0.2 0.1], Bias: -0.1

Iteration no: 3
Input: [0 0], Pred: 0, Error: 0, Updated Weights: [0.2 0.1], Bias: -0.1
Input: [0 1], Pred: 1, Error: -1, Updated Weights: [0.2 0. ], Bias: -0.2
Input: [1 0], Pred: 1, Error: -1, Updated Weights: [0.1 0. ], Bias: -0.30000000000000004
Input: [1 1], Pred: 0, Error: 1, Updated Weights: [0.2 0.1], Bias: -0.20000000000000004

Iteration no: 4
Input: [0 0], Pred: 0, Error: 0, Updated Weights: [0.2 0.1], Bias: -0.20000000000000004
Input: [0 1], Pred: 0, Error: 0, Updated Weights: [0.2 0.1], Bias: -0.20000000000000004
Input: [1 0], Pred: 0, Error: 0, Updated Weights: [0.2 0.1], Bias: -0.20000000000000004
Input: [1 1], Pred: 1, Error: 0, Updated Weights: [0.2 0.1], Bias: -0.20000000000000004

Iteration no: 5
Input: [0 0], Pred: 0, Error: 0, Updated Weights: [0.2 0.1], Bias: -0.20000000000000004
Input: [0 1], Pred: 0, Error: 0, Updated Weights: [0.2 0.1], Bias: -0.20000000000000004
Input: [1 0], Pred: 0, Error: 0, Updated Weights: [0.2 0.1], Bias: -0.20000000000000004
Input: [1 1], Pred: 1, Error: 0, Updated Weights: [0.2 0.1], Bias: -0.20000000000000004

Final AND Predictions:
Input: [0 0], Output: 0
Input: [0 1], Output: 0
Input: [1 0], Output: 0
Input: [1 1], Output: 1

```

## Practical – 12

**Aim: Build an Artificial Neural Network by implementing the Backpropagation Algorithm.**

- **Code:**

```
import numpy as np

def sigmoid(x):
    return 1 / (1 + np.exp(-x))
def sigmoid_derivative(x):
    return x * (1 - x)
def mse(y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)
X = np.array([[0,0], [0,1], [1,0], [1,1]])
y = np.array([[0], [1], [1], [1]])
np.random.seed(42)
input_size = 2
hidden_size = 4
output_size = 1
W1 = np.random.randn(input_size, hidden_size)
b1 = np.zeros((1, hidden_size))
W2 = np.random.randn(hidden_size, output_size)
b2 = np.zeros((1, output_size))
learning_rate = 0.1
epochs = 10000
for epoch in range(epochs):
    # Forward pass
    z1 = np.dot(X, W1) + b1
    a1 = sigmoid(z1)
    z2 = np.dot(a1, W2) + b2
    a2 = sigmoid(z2)
    # Backward pass
    error = y - a2
    d_a2 = error * sigmoid_derivative(a2)
    d_W2 = np.dot(a1.T, d_a2)
    d_b2 = np.sum(d_a2, axis=0, keepdims=True)
    d_a1 = np.dot(d_a2, W2.T) * sigmoid_derivative(a1)
    d_W1 = np.dot(X.T, d_a1)
    d_b1 = np.sum(d_a1, axis=0, keepdims=True)
    # update weights
    W2 += learning_rate * d_W2
    b2 += learning_rate * d_b2
    W1 += learning_rate * d_W1
    b1 += learning_rate * d_b1
    if epoch % 1000 == 0:
        print(f'Epoch {epoch}, Loss: {mse(y, a2):.6f}')
print("\nFinal outputs after training: ")
print(np.round(a2))
```



- **Output:**



```
Epoch 0, Loss: 0.387145  
Epoch 1000, Loss: 0.026790  
Epoch 2000, Loss: 0.005534  
Epoch 3000, Loss: 0.002624  
Epoch 4000, Loss: 0.001640  
Epoch 5000, Loss: 0.001168  
Epoch 6000, Loss: 0.000897  
Epoch 7000, Loss: 0.000723  
Epoch 8000, Loss: 0.000602  
Epoch 9000, Loss: 0.000515
```

Final outputs after training:

```
[[0.]  
 [1.]  
 [1.]  
 [1.]]
```